Dropped columns

I dropped four columns. The first was respondent_email, because it identifies people and the analysis does not need it. The second was ip_address, for the same reason. The third was free_comment_2, which was empty for almost every respondent. The fourth was survey_version, because it had the same value in every response, so it carries no information.

Missing ages

Rows with a missing age were dropped.

Duplicates

Some students submitted the survey twice. I treated two responses as duplicates when they had the same hashed student identifier and the same answers to all the rating questions. When that happened I kept the earlier submission by timestamp and removed the later one. Responses from the same student with different answers were kept, because a changed answer is not a duplicate.

Notebook

All the cleaning steps are in the notebook cleaning.ipynb, in the order described above, and it is meant to be run from the first cell to the last.